



CHALMERS
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF GOTHENBURG

MODULE 3: CLUSTERING 2

DAT405 / DIT407, 2022-2023, READING PERIOD 4


Topics

- DBSCAN clustering
- Hierarchical clustering
- Validating clusterings

Limitations of K-means clustering



K-means: result depends on initialization



DBSCAN clustering

[Steven Bierwagen](#)

DBSCAN

- Density-Based Spatial Clustering of Applications with Noise
- From 1996

Ingredients for DBSCAN

- A *distance* measure (or metric or similarity measure)
 - often Euclidean distance
- A number defining the meaning of *neighbor*
 - epsilon: the max distance between two points considered neighbors.
- A number defining the meaning of *cluster* (vs outlier or noise)
 - minpts: the minimum number of points in a cluster.

scanning
radius

min points
inside radius


Two hyperparameters

Labeling step


All points in dataset labeled as one of these:

- Core point
- Border point
- Noise point

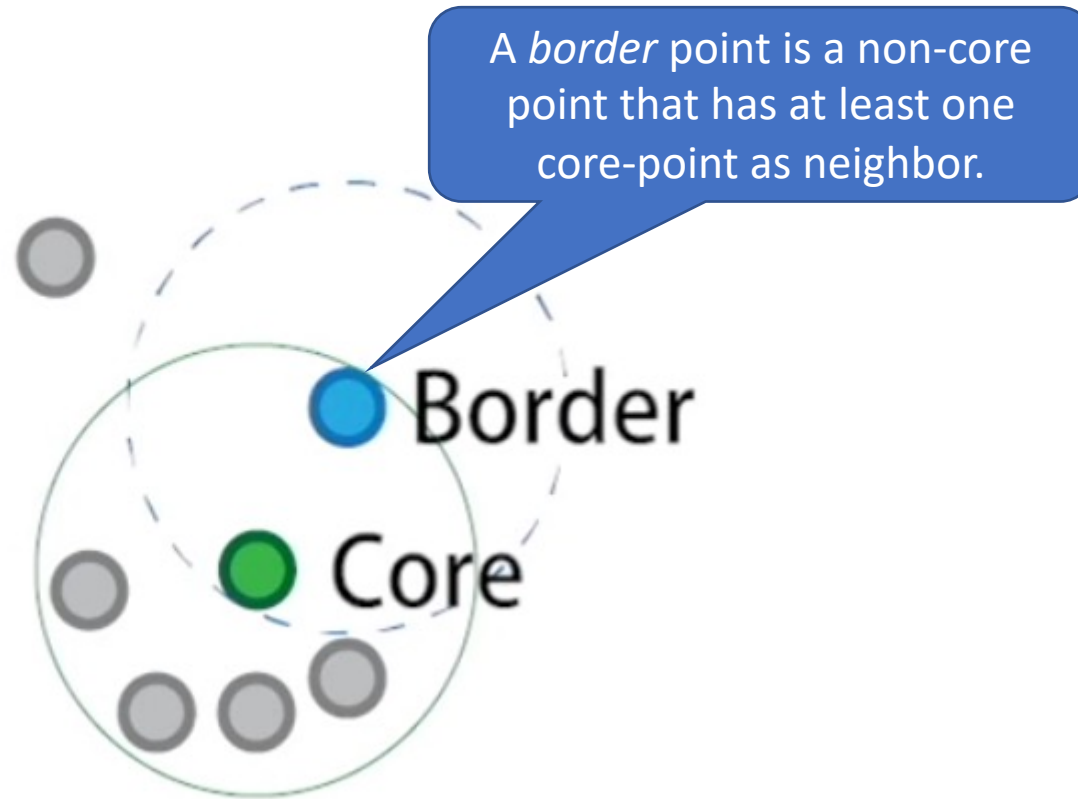
Neighbors




Core points



Border points




Noise points




Clustering step

Clusters all core points and border points. Outliers will not be clustered!


Start by picking a new color c and an uncoloured core point p .



Clustering step



Clustering step






Algorithm


Algorithm 8.4 DBSCAN algorithm.

- 1: Label all points as core, border, or noise points.
 - 2: Eliminate noise points.
 - 3: Put an edge between all core points that are within Eps of each other.
 - 4: Make each group of connected core points into a separate cluster.
 - 5: Assign each border point to one of the clusters of its associated core points.
-

Clusterings created by DBSCAN



Using DBSCAN



See Jupyter notebooks for Module 3 for code examples

K-means vs. DBSCAN

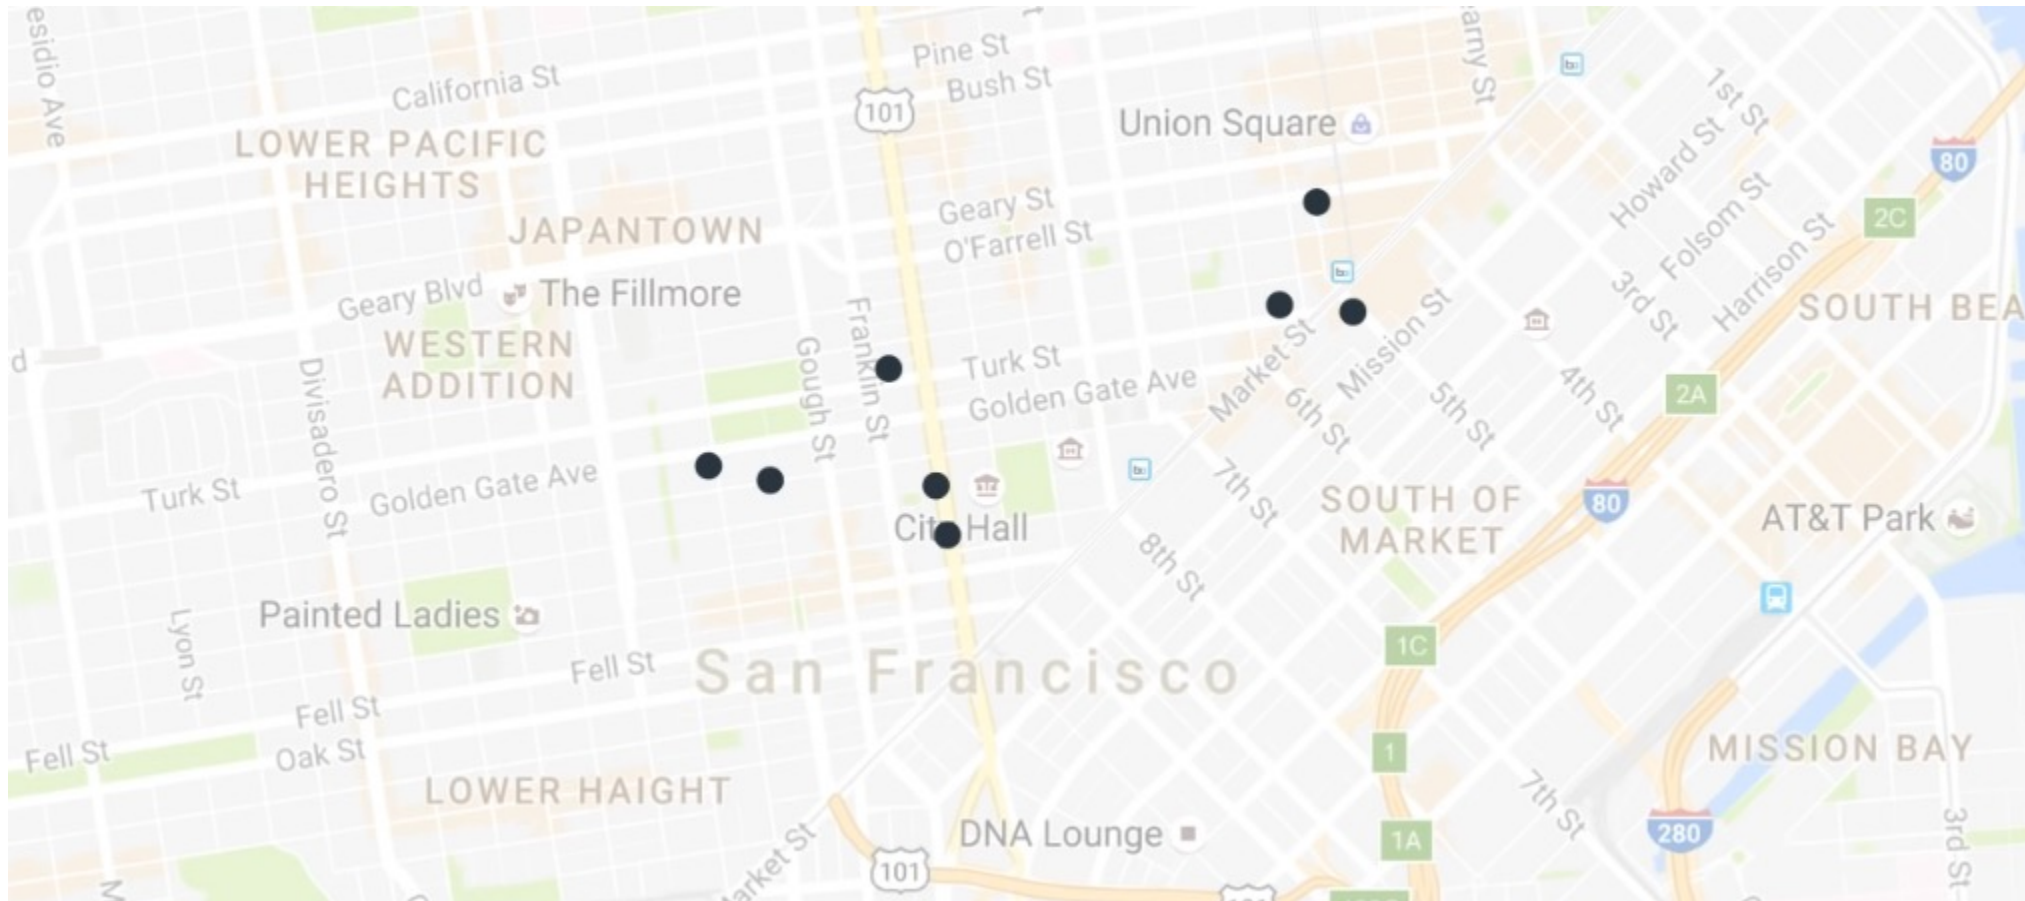
- K-means assigns all points to a cluster, whereas DBSCAN doesn't necessarily do this. DBSCAN treats outliers as outliers.
- K-means works best when clusters are basically spherical. DBSCAN can find arbitrarily-shaped clusters.
- DBSCAN doesn't require the number of clusters to be specified by the user.

Hierarchical clustering

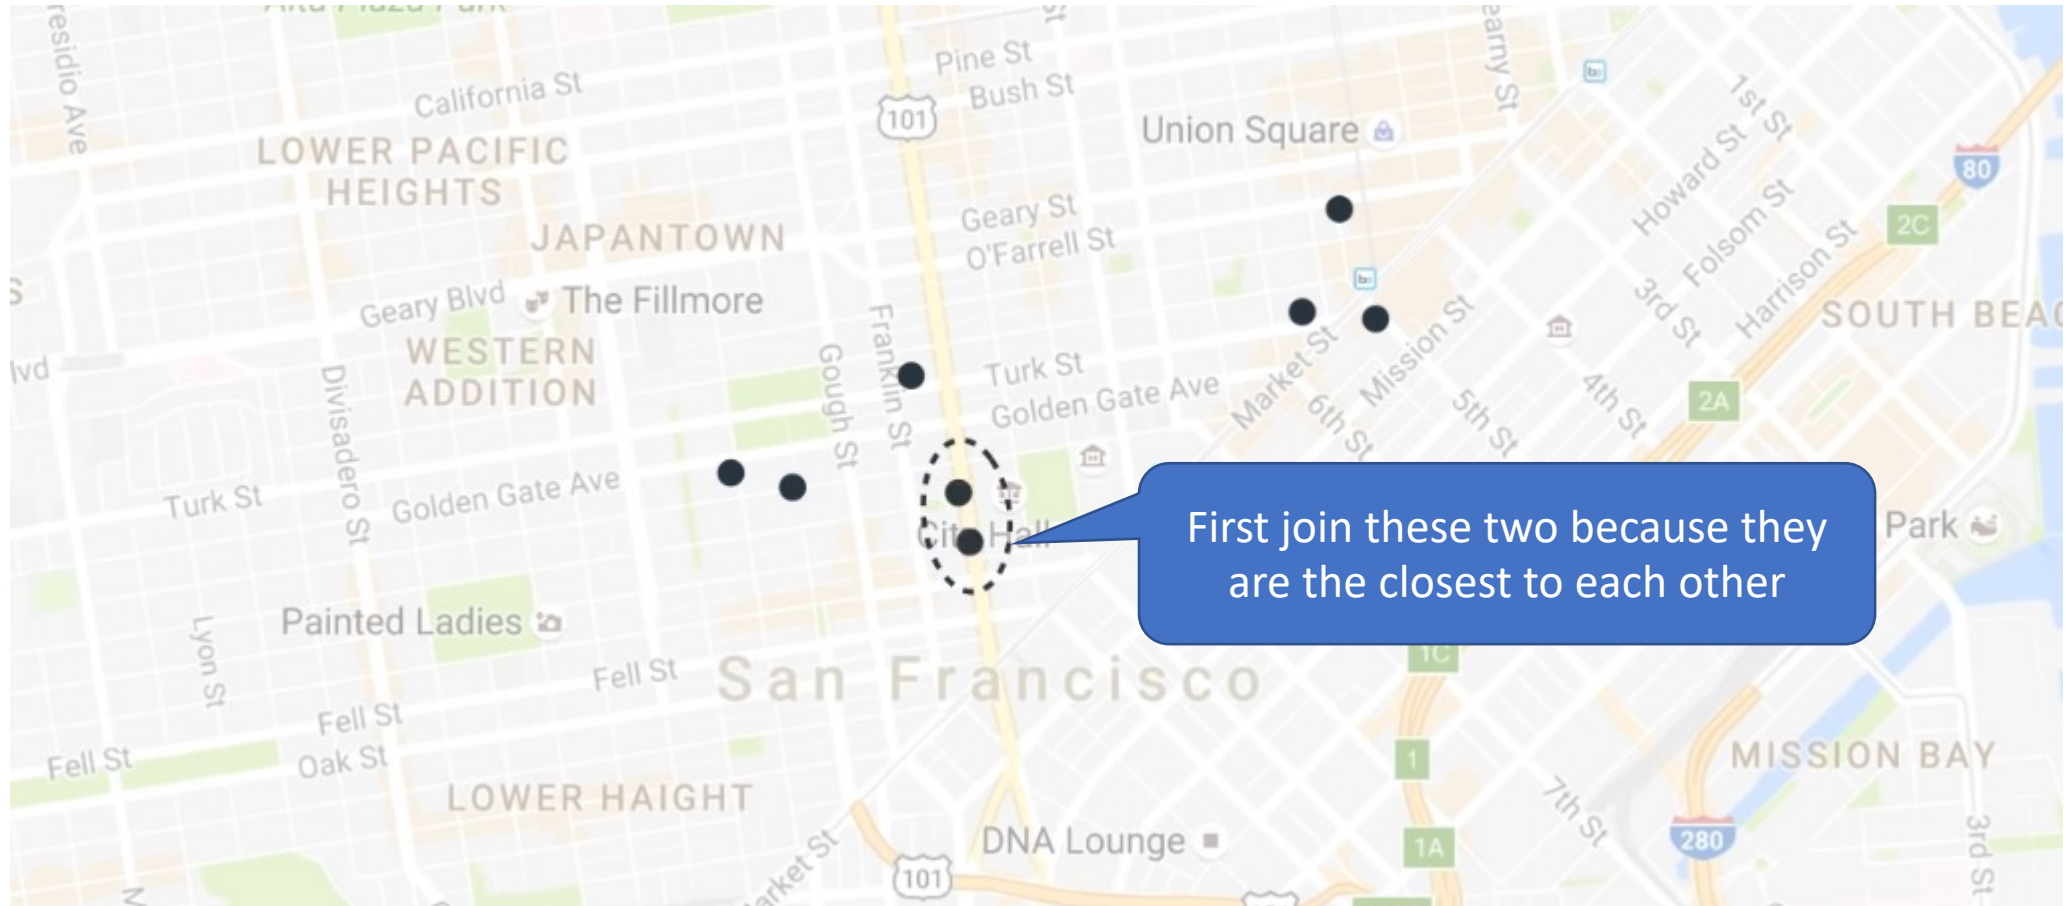
[Luis Serrano](#)

A new set of points

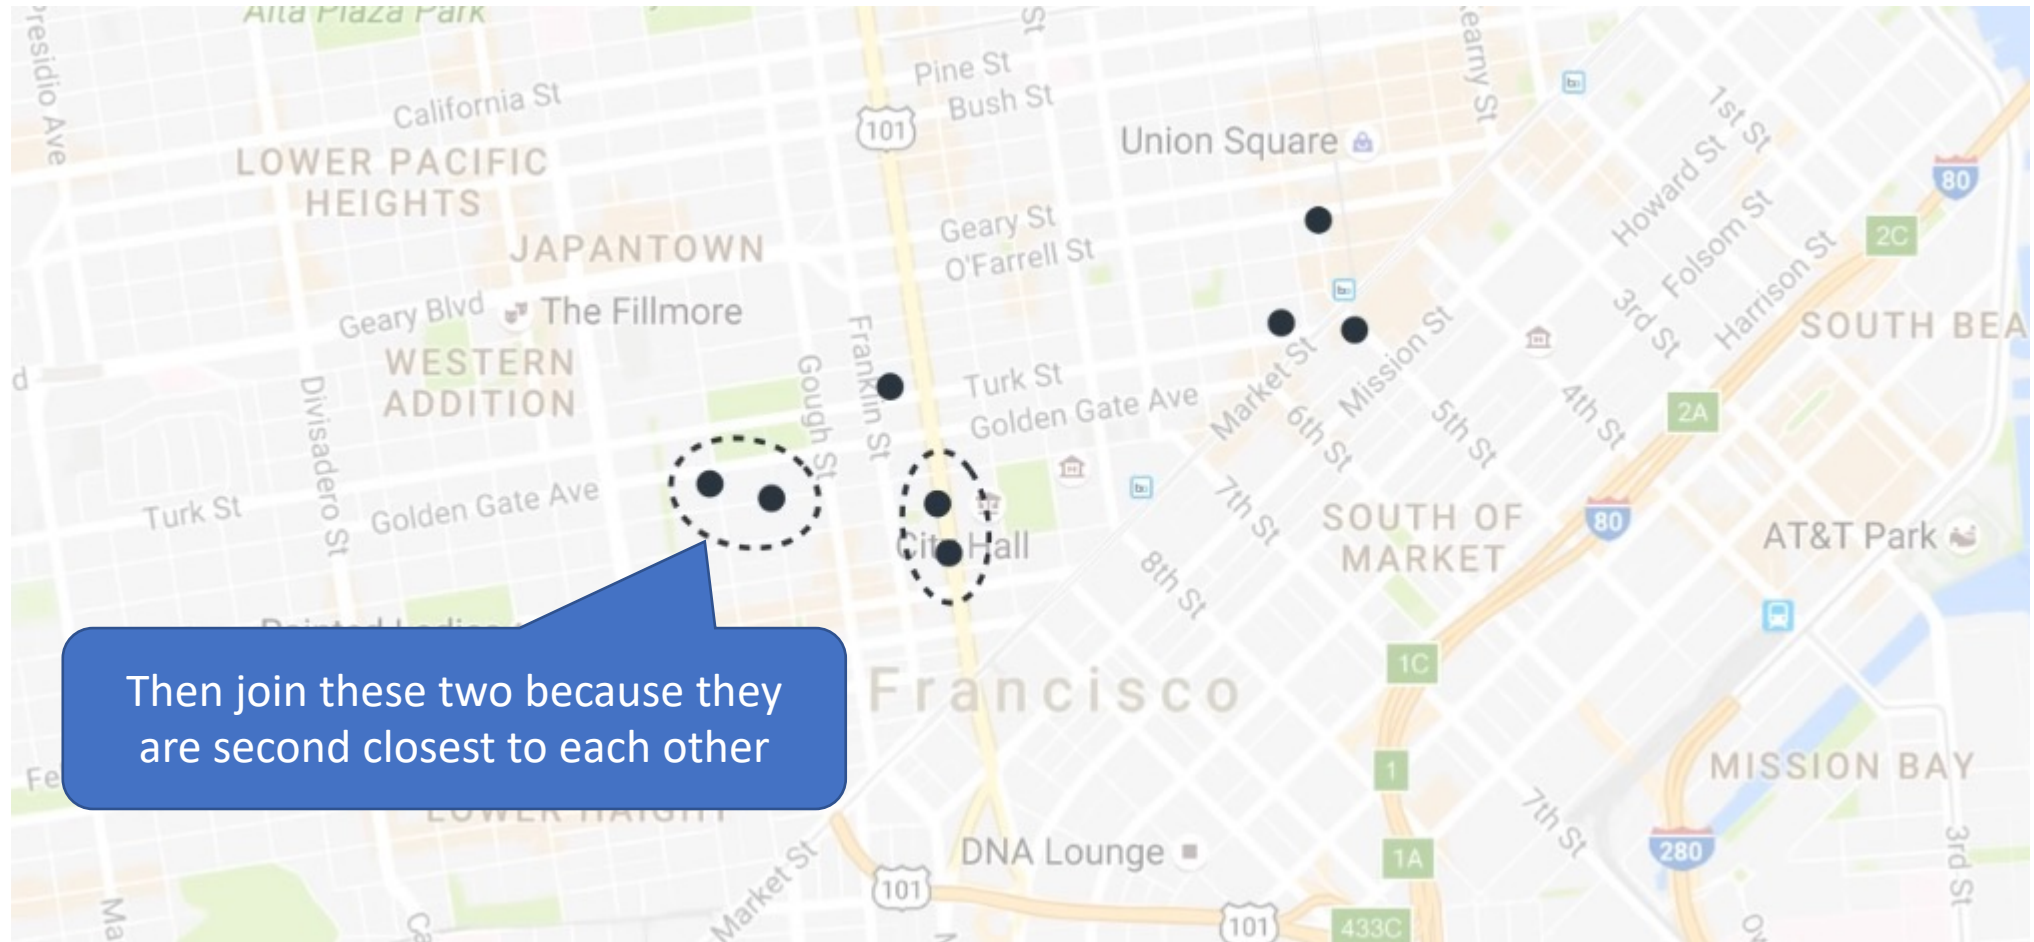
Suppose we want to cluster these addresses by proximity. No pizza parlors involved this time!



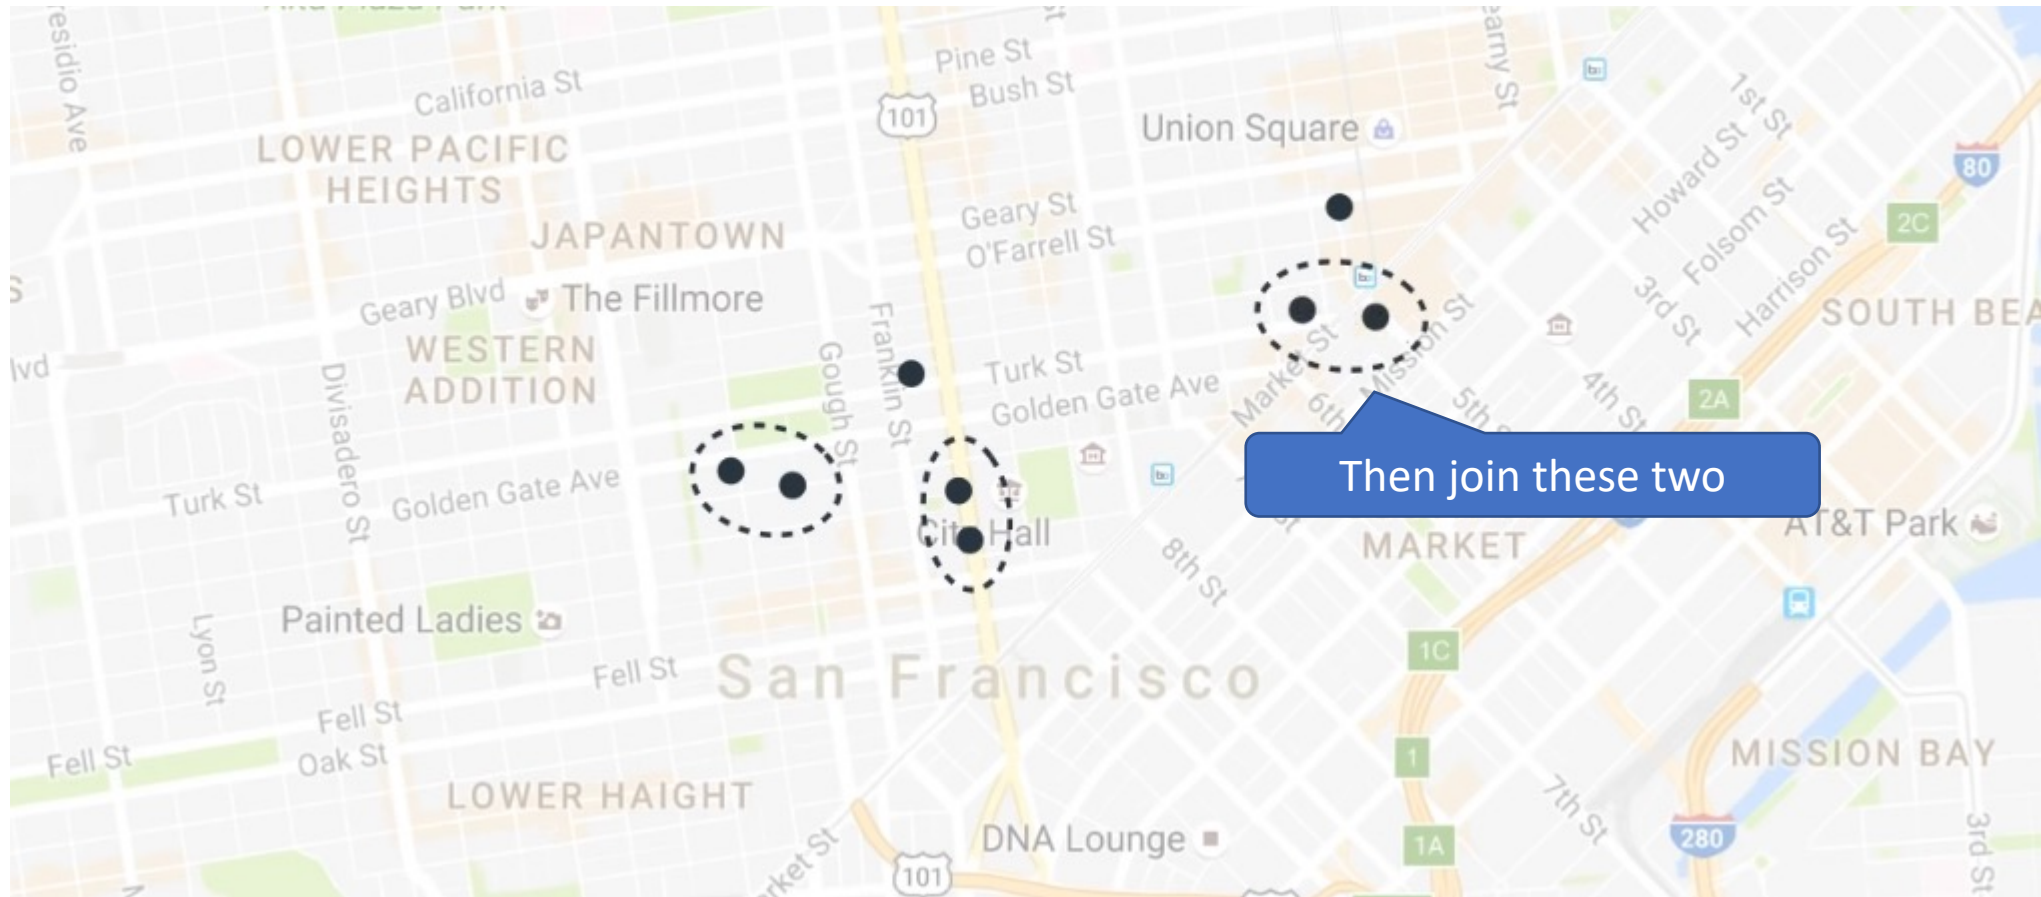
Join the close ones



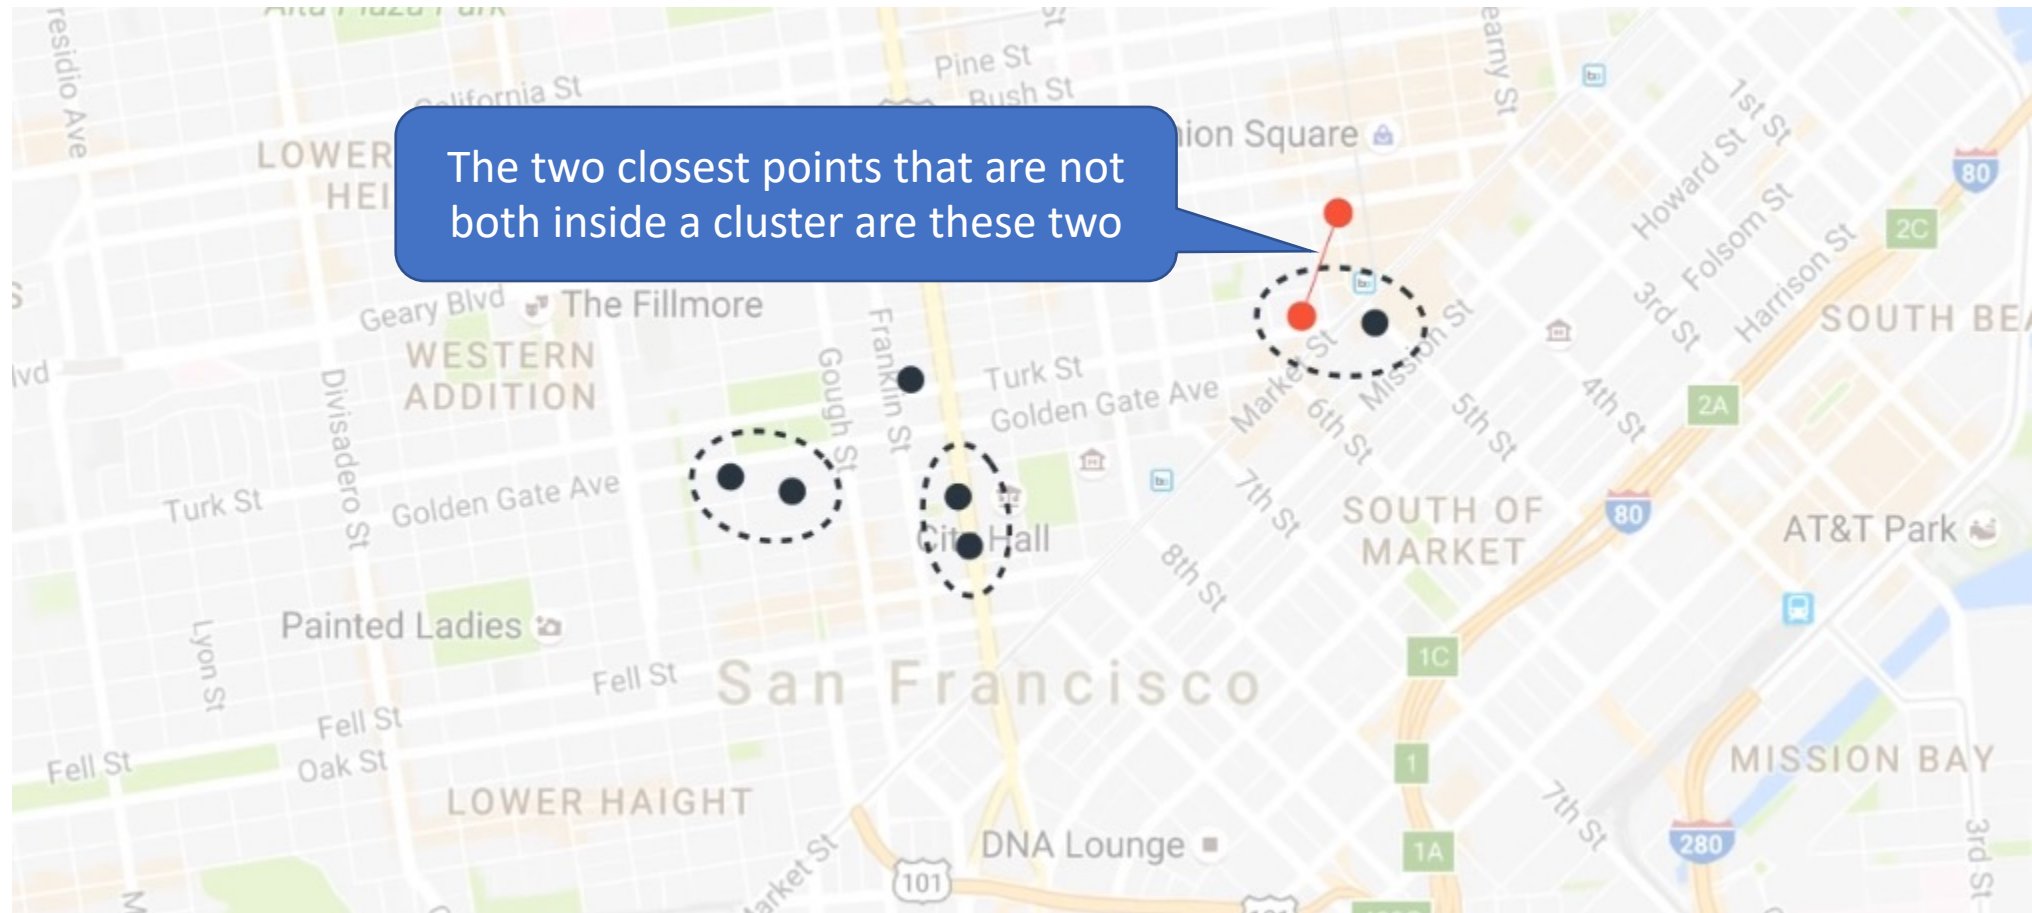
Join the close ones



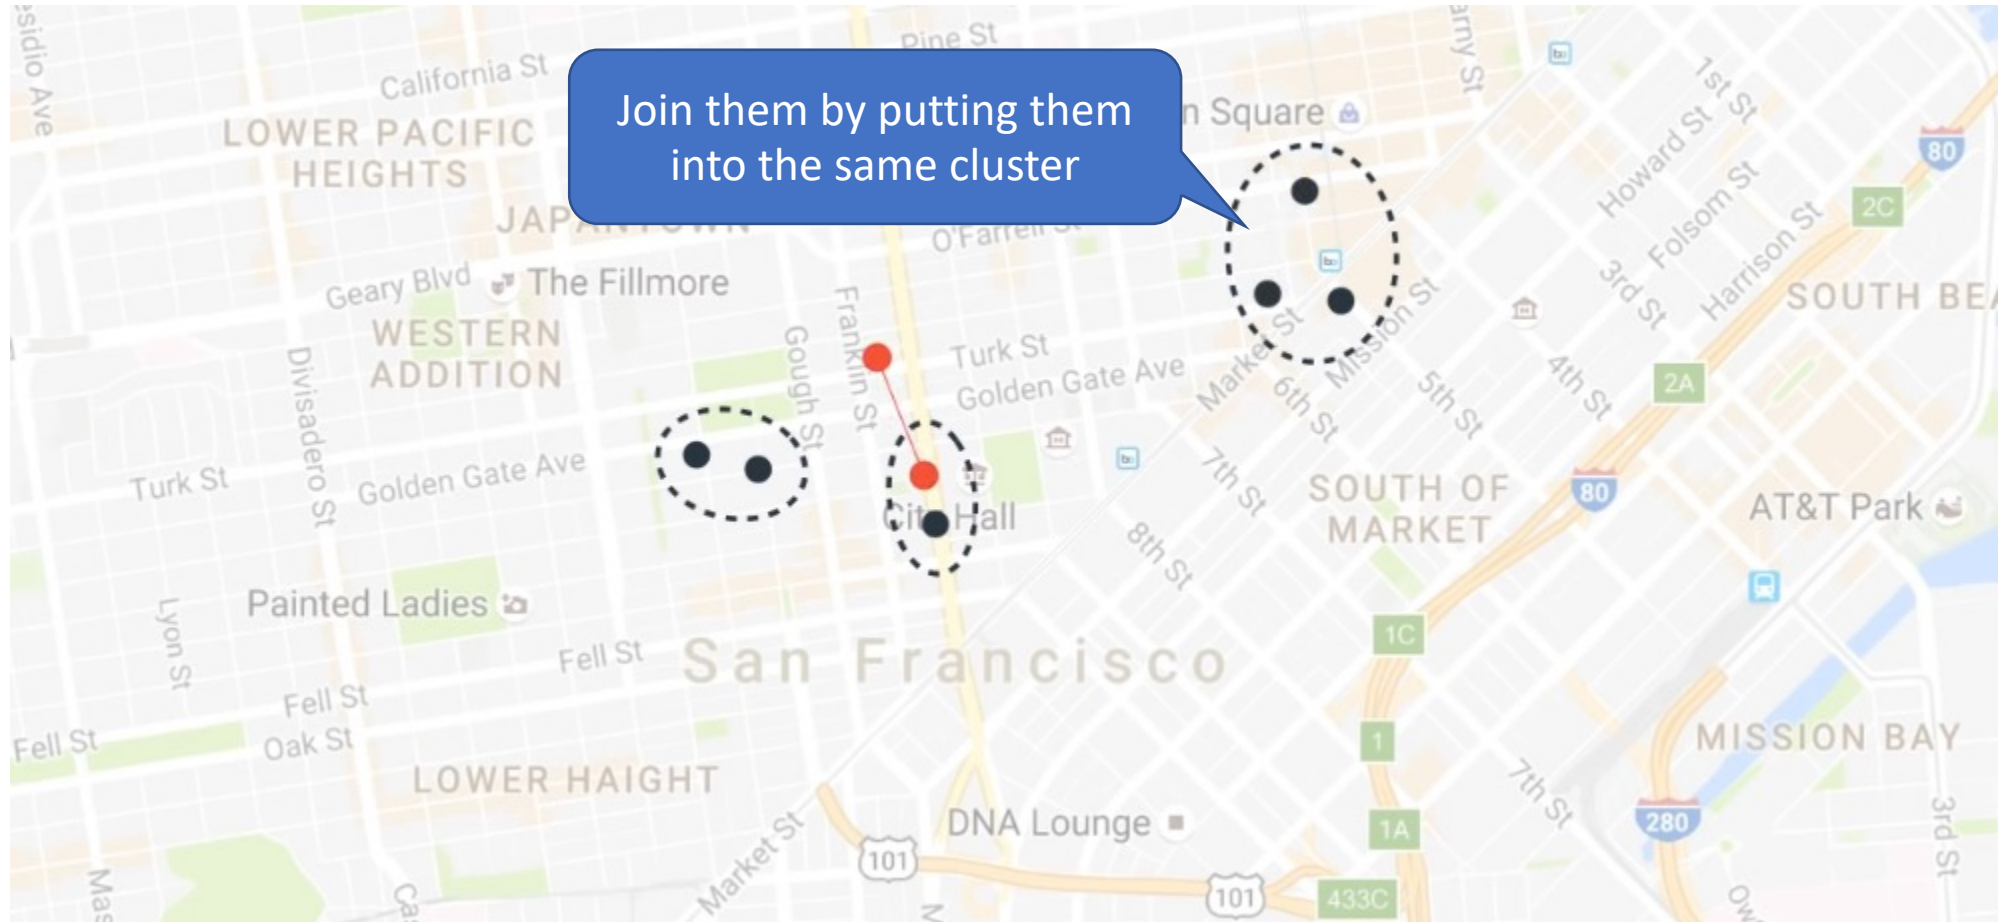
Join the close ones



Join the close ones



Join the close ones



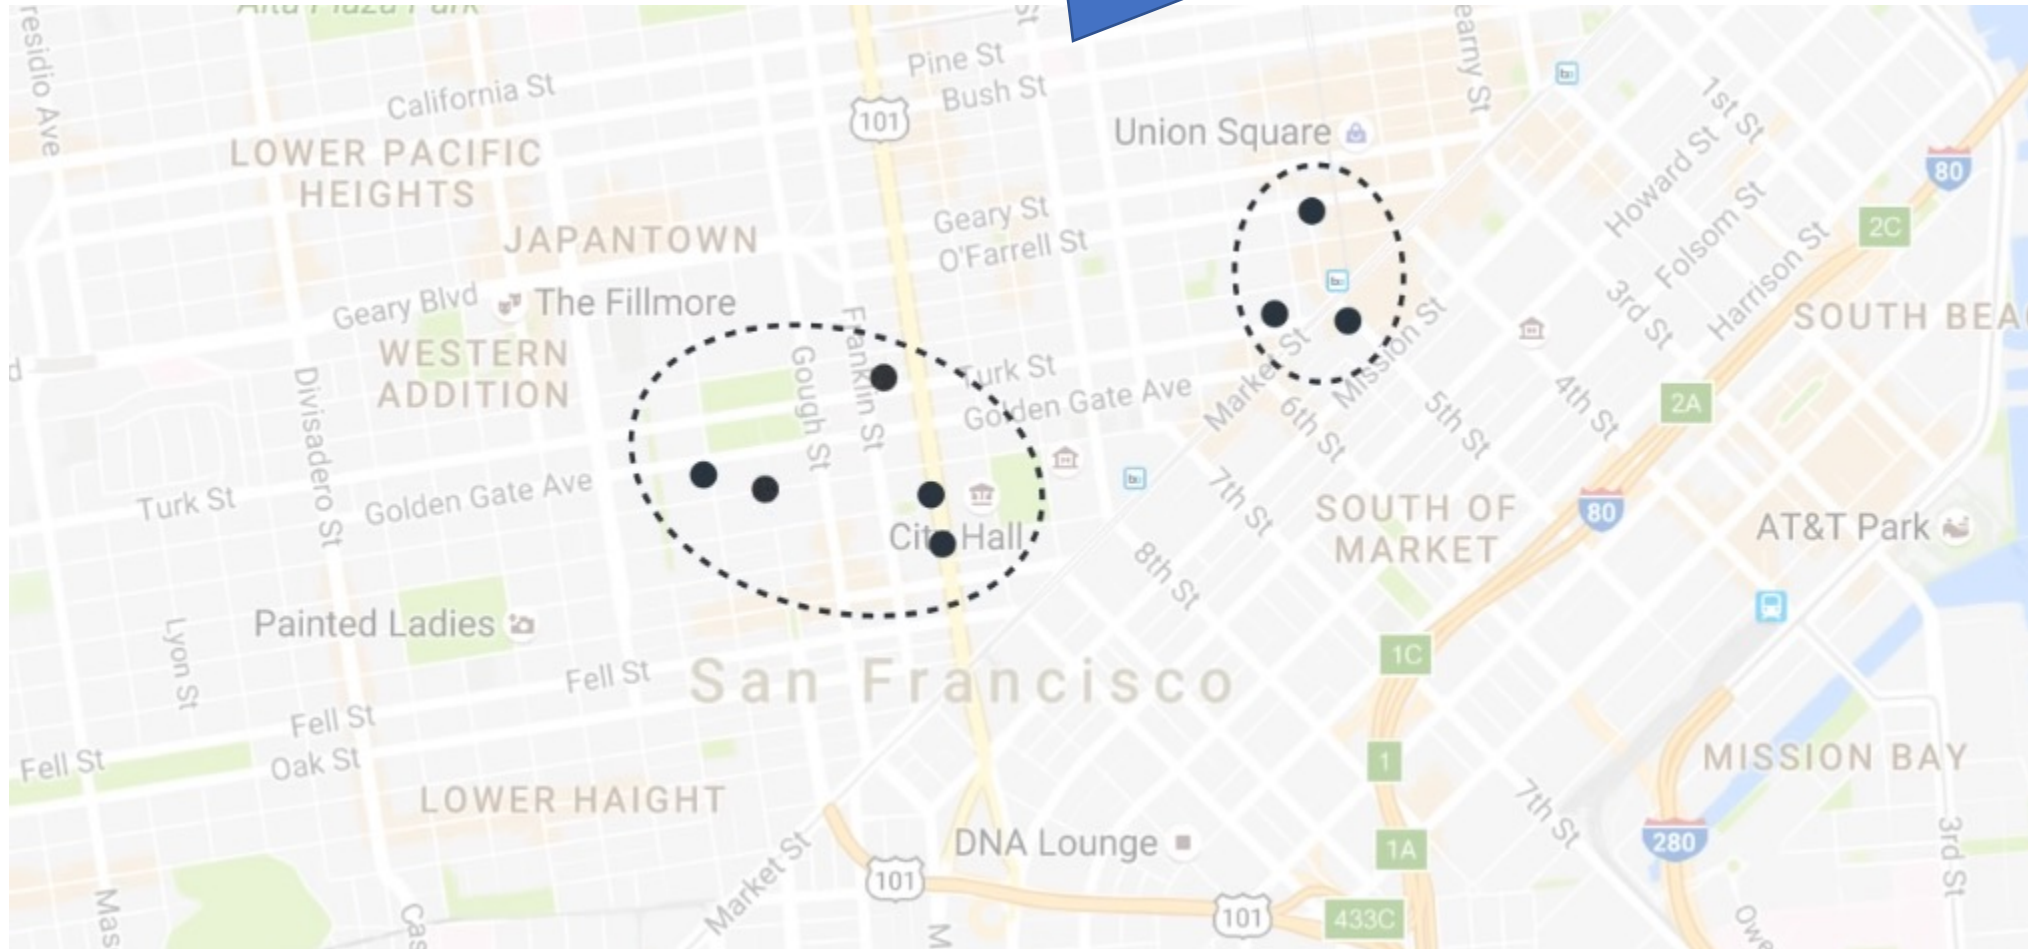
Join the close ones

Continue in the same way...



Join the close ones

Continue in the same way...

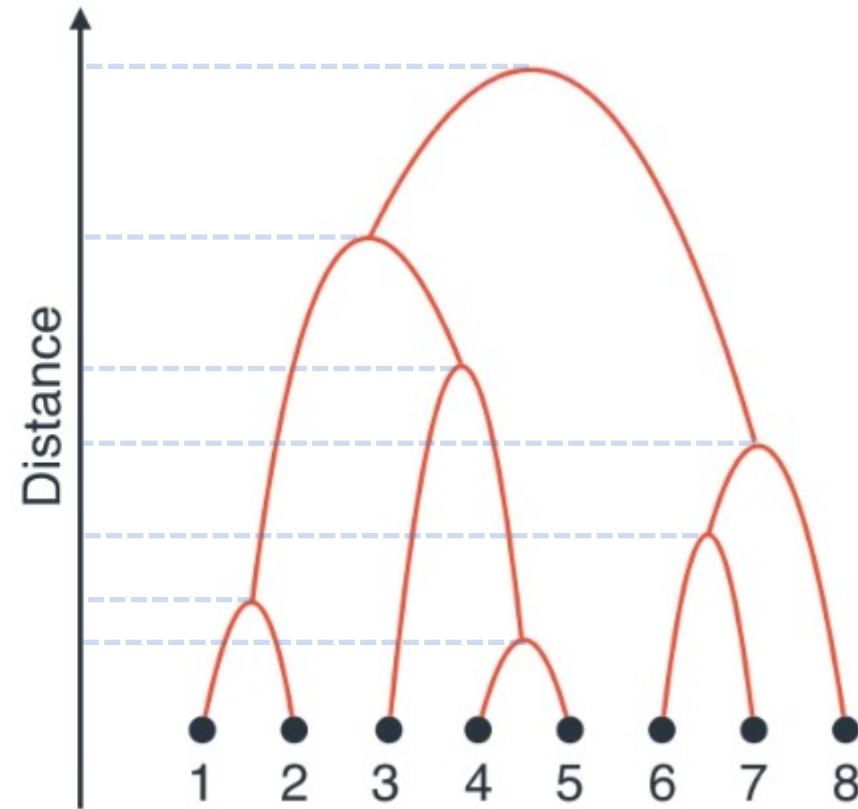



Join the close ones




Dendrogram


A *dendrogram* shows the entire hierarchical clustering process (without STOP)



Dendrogram



Dendrogram




If we have a space with billions of points in thousands of dimensions, the dendrogram is still a 2D graph!

For example the tree of life!

Hierarchical clustering gives more than a clustering: a hierarchy (or taxonomy)

The tree of life



All the major and many of the minor living branches of life are shown on this diagram, but only a few of those that have gone extinct are shown. Example: Dinosaurs - extinct

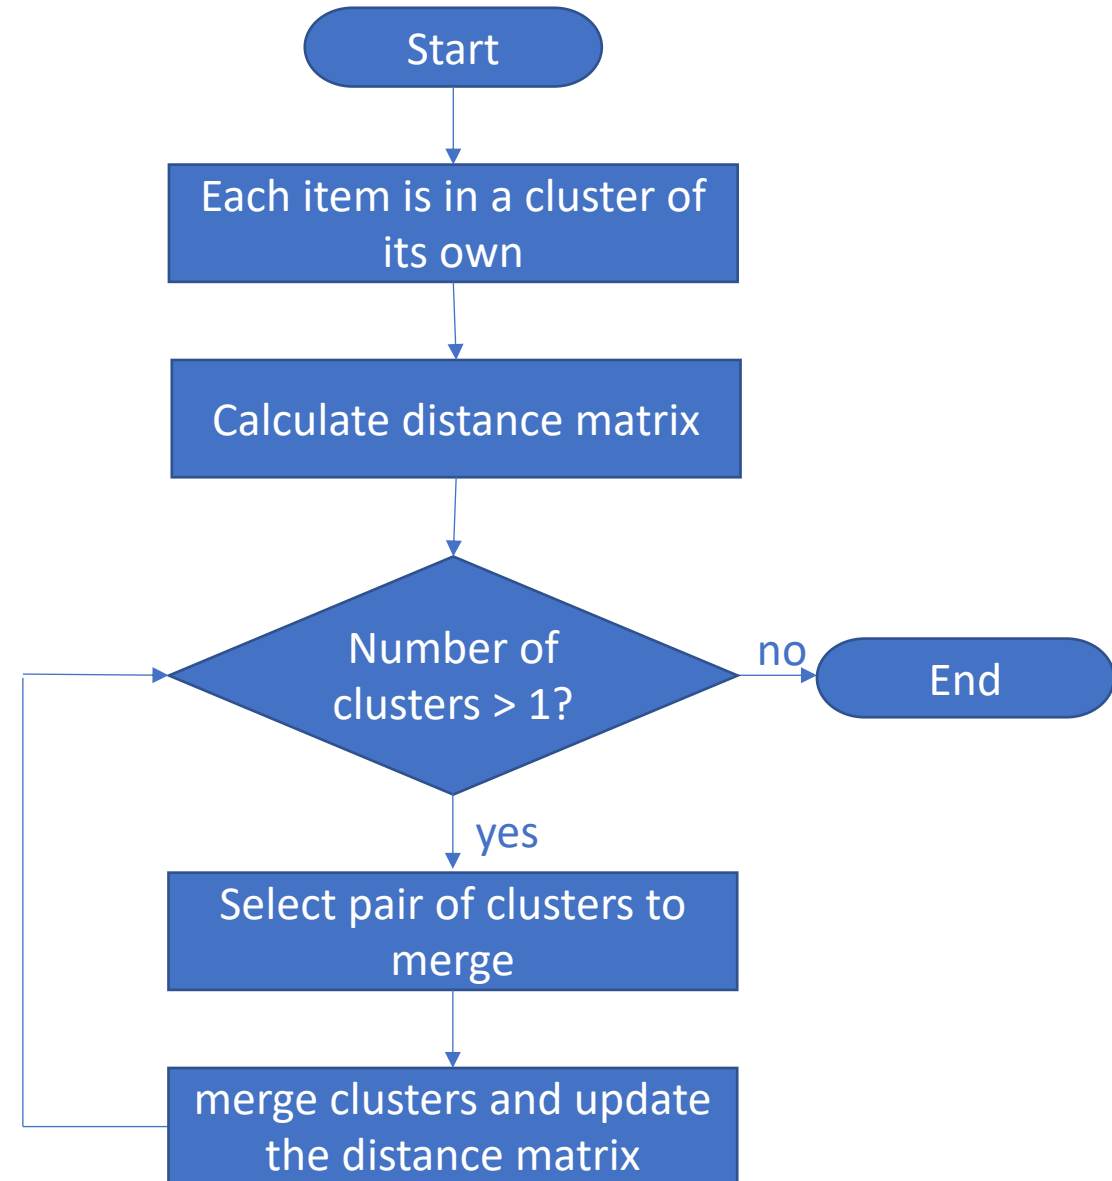


© 2006, 2017 Leonard Eisenberg. All rights reserved.
evogeneao.com

<https://www.evogeneao.com/>

Hierarchical clustering

- Sometimes called agglomerative clustering, when done bottom-up
- From one extreme case (many clusters, each containing one item) to another (one cluster that contains all items)



Distance matrix

Edit distances between protein sequences (strings)

- a. Human haemoglobin alpha chain
- b. Human haemoglobin beta chain
- c. Horse haemoglobin alpha chain
- d. Horse haemoglobin beta chain
- e. Marine bloodworm haemoglobin
- f. Yellow lupine leghaemoglobin

Six proteins with a common evolutionary ancestor

<i>D</i>	a	b	c	d	e	f
a	0	84	18	86	112	121
b	84	0	85	26	117	119
c	18	85	0	84	112	125
d	86	26	84	0	113	121
e	112	117	112	113	0	119
f	121	119	125	121	119	0

Amino acid sequences of six proteins

> human_alpha

VLSPADKTNVKAAGKVGAGHAGEYGAEALERMFLSFPTTKTYFPHFDLSHGSAQVKGHGKKVADALTNAVAHVDDMPNALSALSDLHAHKLRVDPVNFKLLSHCLLVTLAAHLPAEFTPAVHASLDKFLASVSTVLTISKYR

> human_beta

VHLTPEEKSAVTALWGKVNVDEVGGEALGRLLVVYPWTQRFFESFGDLSTPDAMGNPKVKAHGKKVLGAFSDGLAHLNLRGTATLSELHCDKLHVDPENFRLLGNVLCVLAHFGKEFTPPVQAAYQKVVAGVANALAHKYH

> horse_alpha

VLSAADKTNVKAWSKVGGHAGEYGAEALERMFLGFPTTKTYFPHFDLSHGSAQVKAHGKKVADGLTLAVGHLLDLPGLSDLSNLHAHKLRVDPVNFKLLSHCLLVTLAVHLPNDFTPAVHASLDKFLSSVSTVLTISKYR

> horse_beta

VQLSGEEKAAVLALWDKVNNEEVGGEALGRLLVVYPWTQRFFDSFGDLSPGAVMGPNPKVKAHGKKVLHSFGEGVHHLNLRGTFAALSELHCDKLHVDPENFRLLGNVLAIVARHFGKDFTPPELQASYQKVVAGVANALAHKYH

> marine_bloodworm

GLSAAQRQVIAATWKDIAGADNGAGVGKKCLIKFLSAHPQMAAVFGFSGASDPGVAALGAKVLAQIGVAVSHLGDEGKMVAQMKAVGVRHKGYNKHIKAQYFEPLGASLLSAMEHRIGGKMNAAKDAWAAAYADISGALISGLQS

> yellow_lupine

GALTESQAALVKSSWEEFNANIPKHTRFFILVLEIAPAAKDLFSFLKGTSEVPQNNPELQAHAGKVFKLVYEAAIQLEVTGVVVTDATLKNLGSVHVSKGVADAHFPVVKAEILKTIKEVVGAKWSEELNSAWTIAYDELAIVIKKEMDDAA

Edit distance is the number of single character operations that are required to change one string into another.

Merging clusters

- When clusters **u** and **v** are merged, how do we calculate the distance between the merged cluster and each of the other clusters?
- Various algorithms to choose from, e.g.
 - complete linkage (furthest inter-cluster distance) $\max(\text{dist}(u[i]), v[j]))$
 - single linkage (closest inter-cluster distance) $\min(\text{dist}(u[i]), v[j]))$
 - average linkage
 - Unweighted Pair Group Method with Arithmetic Mean (UPGMA)
 - Weighted Pair Group Method with Arithmetic Mean (WPGMA)
 - ... and many more

See e.g.

<https://docs.scipy.org/doc/scipy/reference/generated/scipy.cluster.hierarchy.linkage.html#scipy.cluster.hierarchy.linkage>

Example: merging clusters

<i>D</i>	a	b	c	d	e	f
a	0	84	18	86	112	121
b		0	85	26	117	119
c			0	84	112	125
d				0	113	121
e					0	119
f						0

1) Shortest distance a – c

2) Merge {a,c}

3) Recompute distance matrix, use max distance (complete linkage)

<i>D</i>	a,c	b	d	e	f
a,c	0	85	86	112	125
b		0	26	117	119
d			0	113	121
e				0	119
f					0

- a. Human haemoglobin alpha chain
- b. Human haemoglobin beta chain
- c. Horse haemoglobin alpha chain
- d. Horse haemoglobin beta chain
- e. Marine bloodworm haemoglobin
- f. Yellow lupine leghaemoglobin

Example: merging clusters

<i>D</i>	a	b	c	d	e	f
a	0	84	18	86	112	121
b		0	85	26	117	119
c			0	84	112	125
d				0	113	121
e					0	119
f						0

- 1) Shortest distance a – c
- 2) Merge {a,c}
- 3) Recompute distance matrix, use max distance between points

<i>D</i>	a,c	b	d	e	f
a,c	0	85	86	112	121
b		0	26	117	119
d			0	113	121
e				0	119
f					0

- 1) Shortest distance b – d
- 2) Merge {b,d}
- 3) Recompute distance matrix

- a. Human haemoglobin alpha chain
- b. Human haemoglobin beta chain
- c. Horse haemoglobin alpha chain
- d. Horse haemoglobin beta chain
- e. Marine bloodworm haemoglobin
- f. Yellow lupine leghaemoglobin

Example: merging clusters

<i>D</i>	a	b	c	d	e	f
a	0	84	18	86	112	121
b		0	85	26	117	119
c			0	84	112	125
d				0	113	121
e					0	119
f						0

- 1) Shortest distance a – c
- 2) Merge {a,c}
- 3) Recompute distance matrix, use max distance between points

<i>D</i>	a,c	b	d	e	f
a,c	0	85	86	112	121
b		0	26	117	119
d			0	113	121
e				0	119
f					0

- 1) Shortest distance b – d
- 2) Merge {b,d}
- 3) Recompute distance matrix


<i>D</i>	a,c	b,d	e	f
a,c	0	86	112	121
b,d		0	117	121
e			0	119
f				0

- a. Human haemoglobin alpha chain
- b. Human haemoglobin beta chain
- c. Horse haemoglobin alpha chain
- d. Horse haemoglobin beta chain
- e. Marine bloodworm haemoglobin
- f. Yellow lupine leghaemoglobin

Validating clustering


Stability on subsets

Clustering stable if removing a proportion of random points does not change the clustering fundamentally




Stability on subsets

Note colors change as labeling clusters into first, second, third ... changes!




Co-occurrence

For all pairs (i,j) count how frequently i and j are in the same cluster.



Co-occurrence




Silhouette coefficient

a: The mean distance between a sample and all other points in the same class.

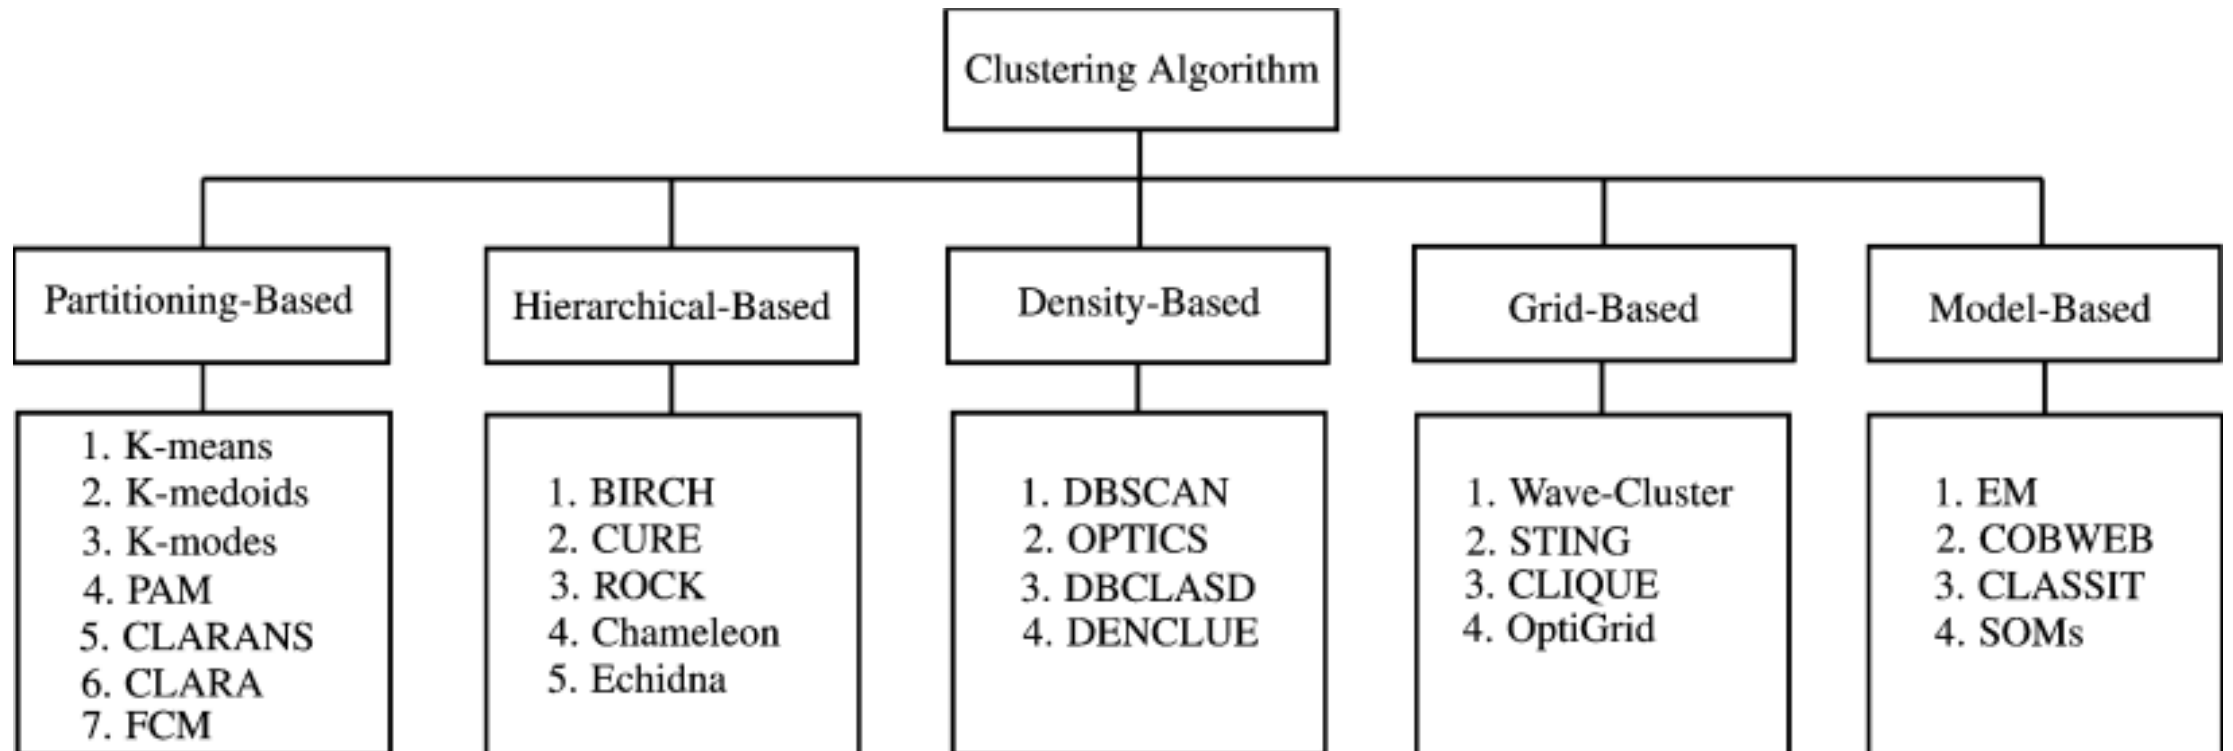
b: The mean distance between a sample and all other points in the *next nearest cluster*.

$$s = \frac{b - a}{\max(a, b)}$$



Ranges between -1 and 1. High value indicate good separation between clusters.

Clustering clustering algorithms



Useful idea when
labeling is
expensive

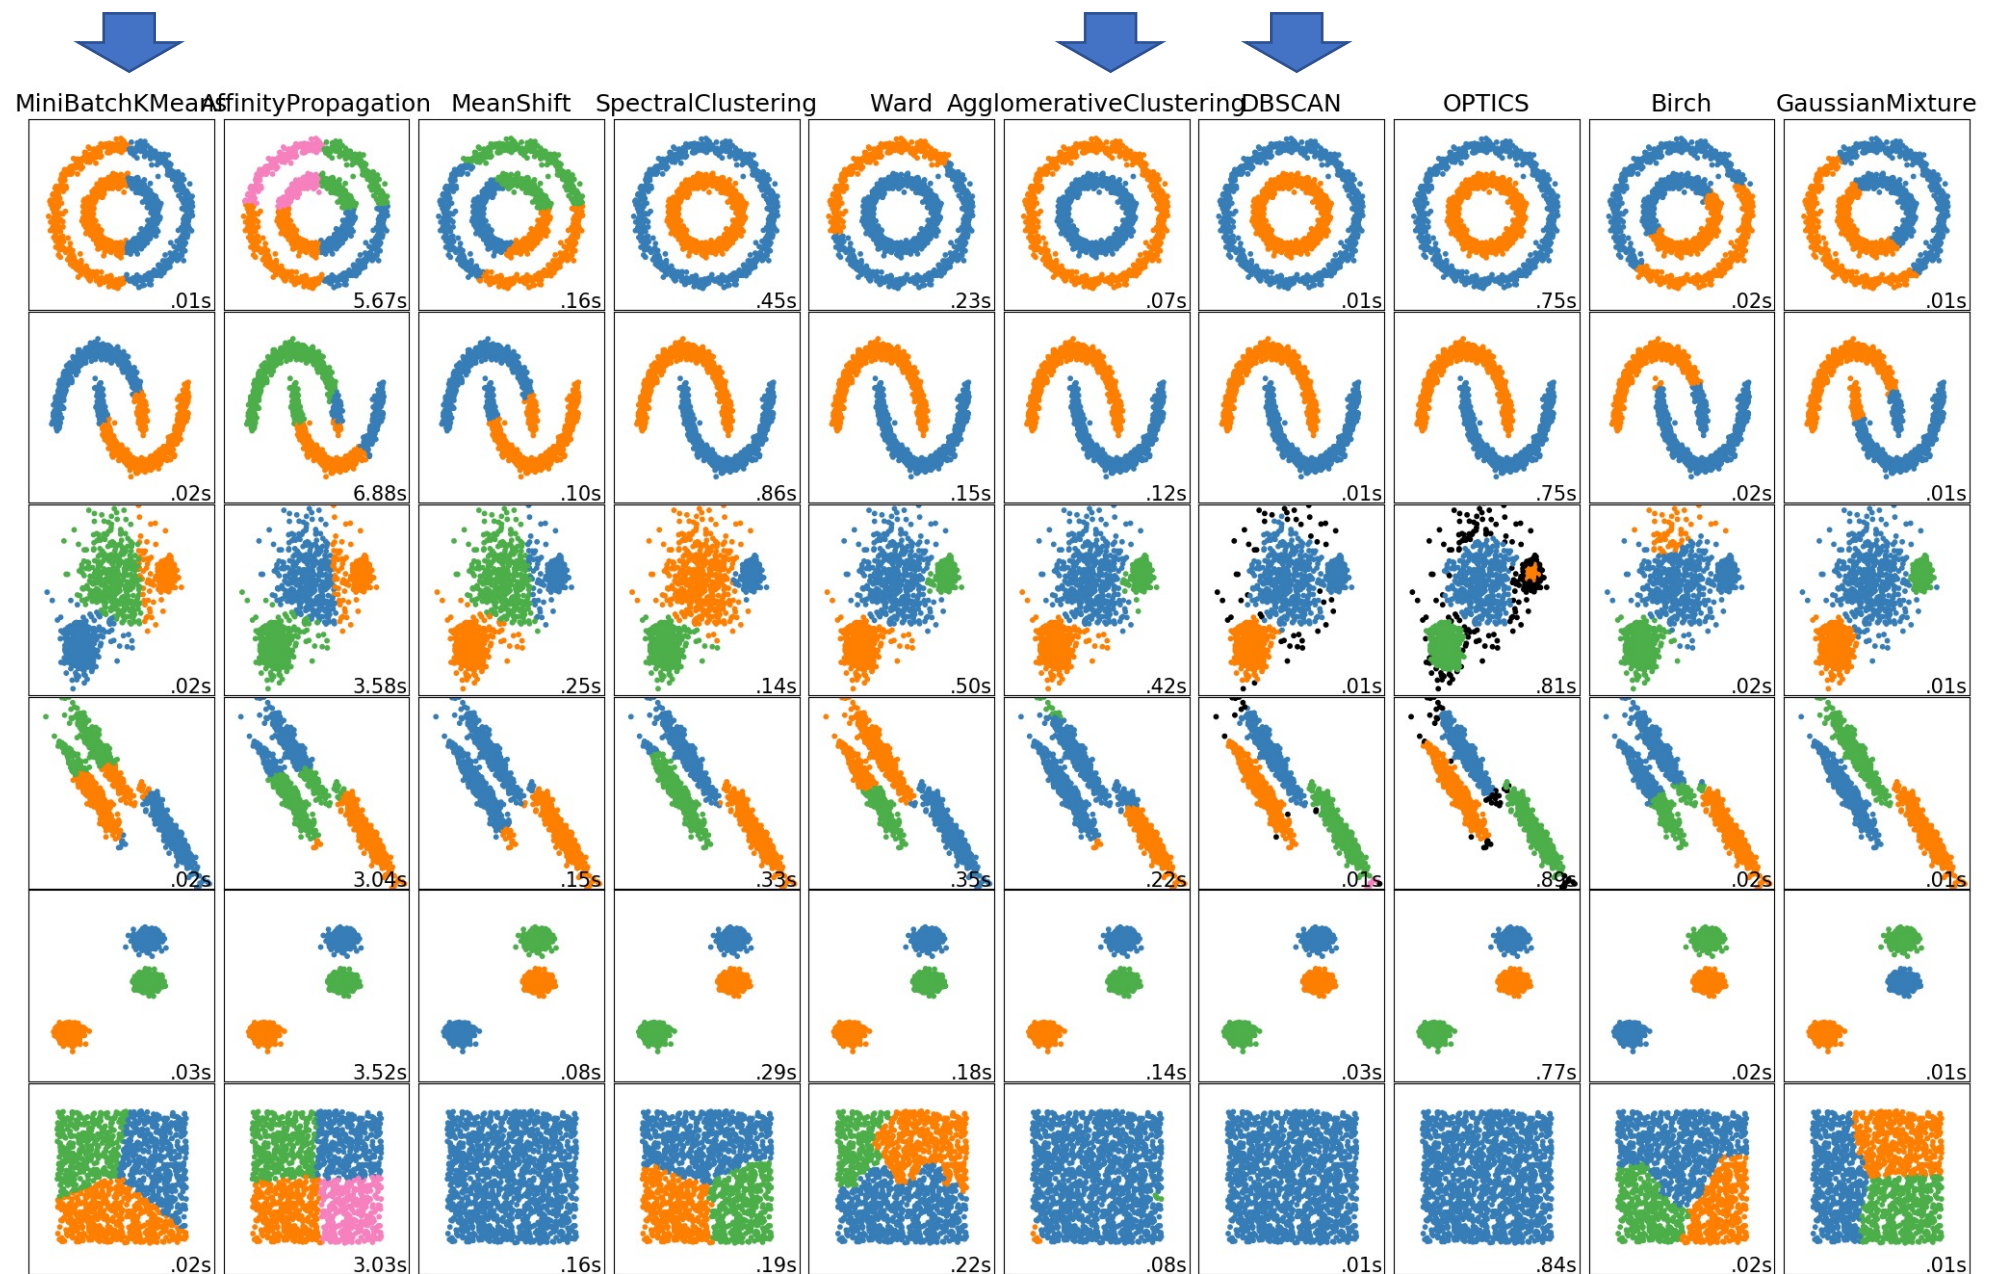
Combining clustering and classification

- Take a dataset with handwritten digits
- Provide only one label per digit (10 labels for the whole dataset)
- Use 10-means with the ten labeled images as starting points for clustering the whole dataset.
- Then use 1nn for classifying new handwritten digits.

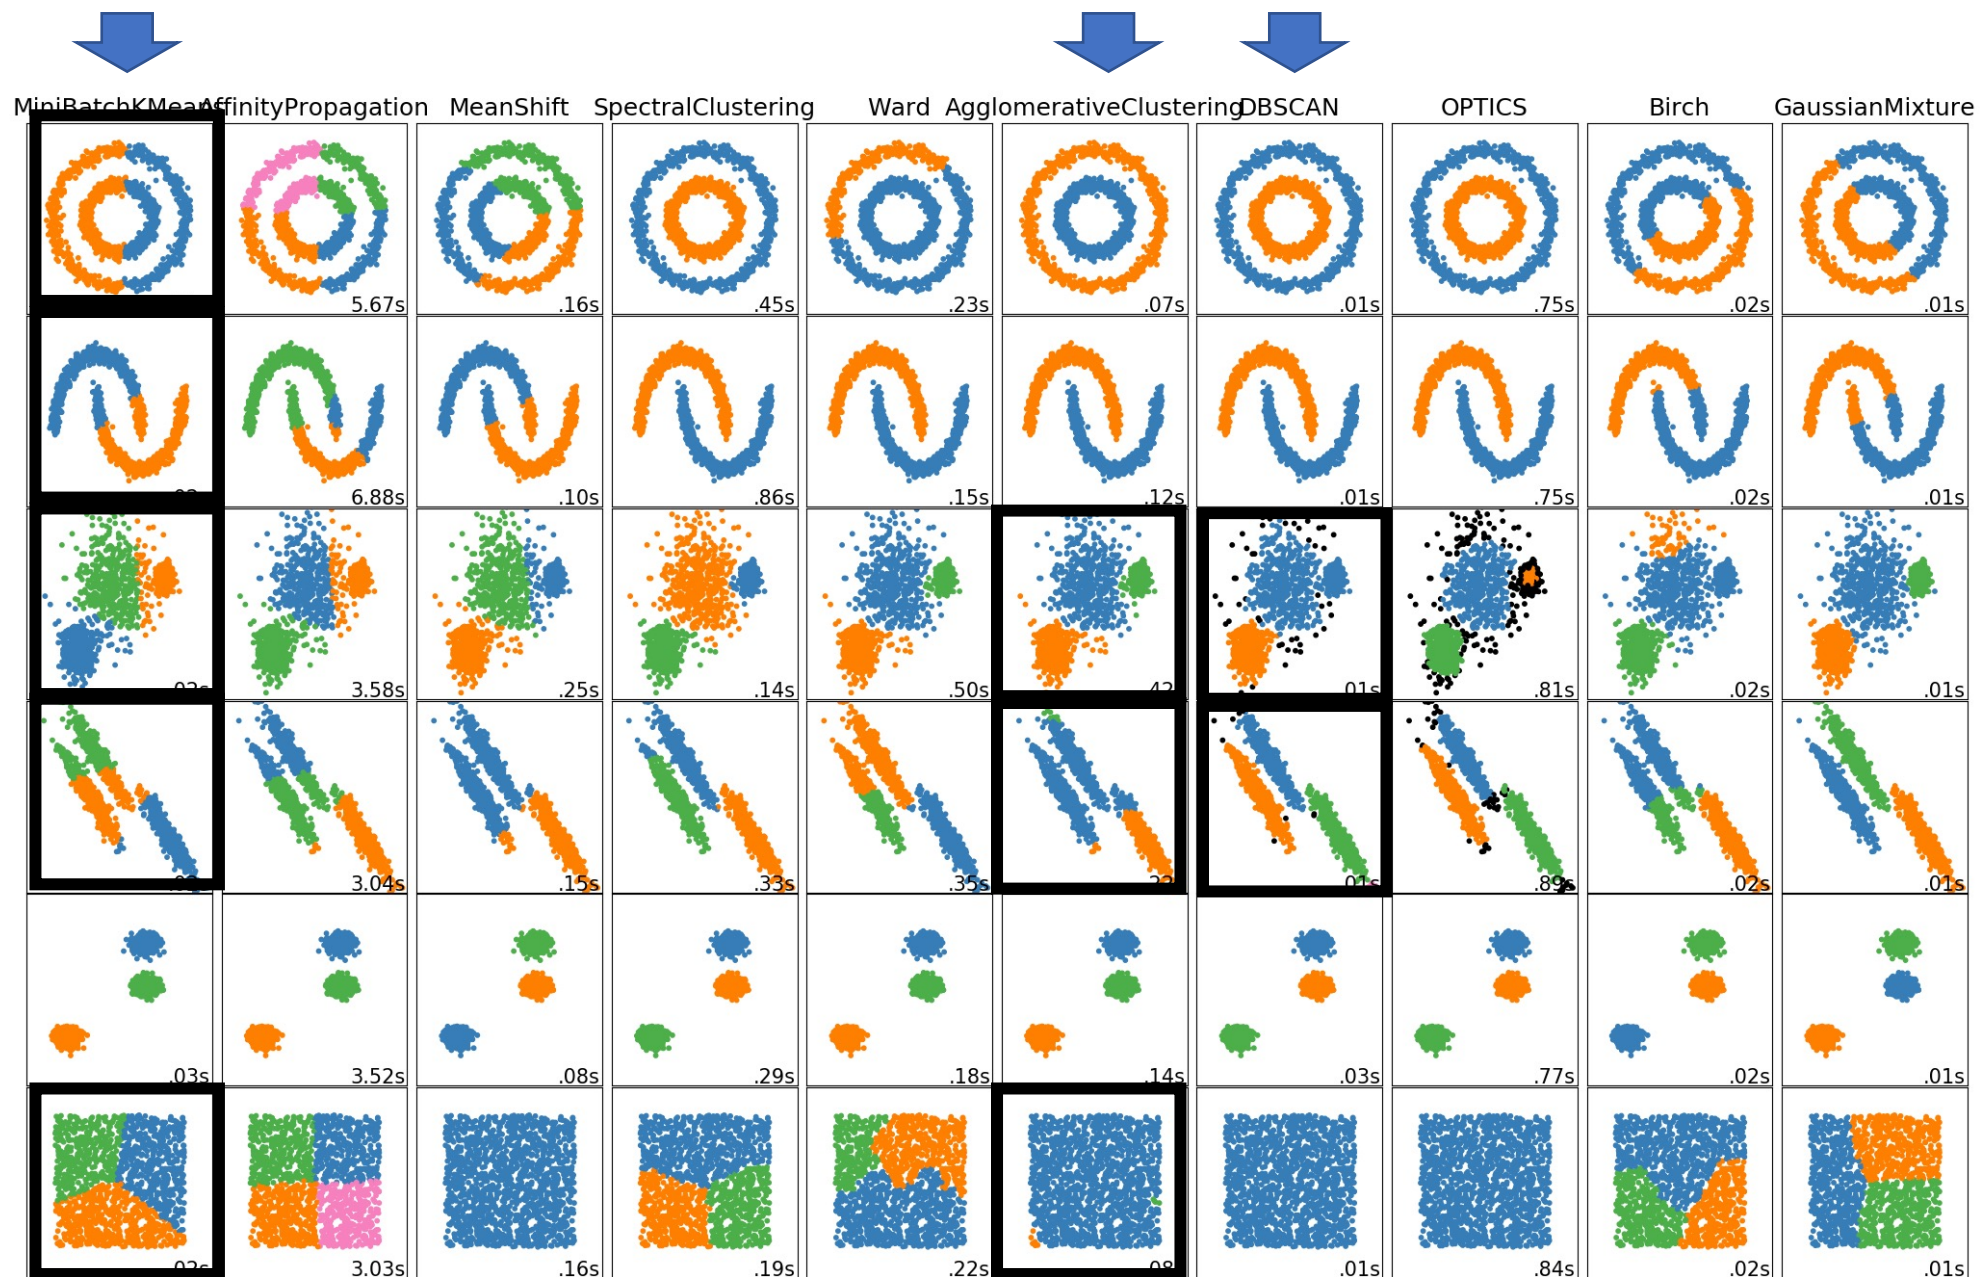
Reflections on clustering

Clustering is successful, but difficult

- Inherent vagueness in the definition of a cluster
- Can be difficult to define an appropriate similarity measure



https://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html




Are the framed cases as desired?


Assignment 3

- Using K-means and density-based clustering to cluster the main chain conformations of amino acid residues in proteins.
- If curious for more information on the problem domain, look at:
 - <http://bioinformatics.org/molvis/phipsi/>
 - <http://tinyurl.com/RamachandranPrincipleYouTube>

Protein main chain



Ramachandran plot



Around 100000
data points
shown here